

# poi\_id

November 14, 2015

```
In [12]: #!/usr/bin/python
```

```
import sys
import pickle
sys.path.append("../tools/")

from feature_format import featureFormat, targetFeatureSplit
from tester import test_classifier, dump_classifier_and_data
%matplotlib inline
import matplotlib.pyplot
from time import time

from sklearn.feature_selection import SelectKBest
from sklearn.feature_selection import f_classif

from sklearn.preprocessing import MinMaxScaler

from sklearn.decomposition import PCA

from sklearn.naive_bayes import GaussianNB
from sklearn.linear_model import LogisticRegression
from sklearn.svm import SVC
from sklearn.tree import DecisionTreeClassifier
from sklearn.ensemble import RandomForestClassifier

from sklearn.pipeline import Pipeline

from sklearn.cross_validation import StratifiedShuffleSplit

from sklearn.grid_search import GridSearchCV
```

```
In [13]: ### Task 1: Select what features you'll use.
### features_list is a list of strings, each of which is a feature name.
### The first feature must be "poi".
features_list = ['poi', 'salary'] # You will need to use more features

### Load the dictionary containing the dataset
data_dict = pickle.load(open("final_project_dataset.pkl", "r") )

# The data_dict is organized as a list of people, each containing a dictionary of features
# Display general information about the dataset...
print "Total Number of data points:", len(data_dict)
print data_dict.keys(), "\n"
```




```

In [14]: ### Task 2: Remove outliers
         #
         # Identify possible outliers by looking at a histogram for each field.
         # Identify the person with the max value for each field.


for feature in data_dict[data_dict.keys()[0]].keys():
    if feature == "email_address" or feature == 'poi':
        continue
    data = featureFormat(data_dict, [feature])
    for person in data_dict:
        if data_dict[person][feature] == max(data):
            print "Possible Outlier for", feature, ":", person
    matplotlib.pyplot.hist(data, bins=10)
    matplotlib.pyplot.xlabel(feature)
    matplotlib.pyplot.show()

```

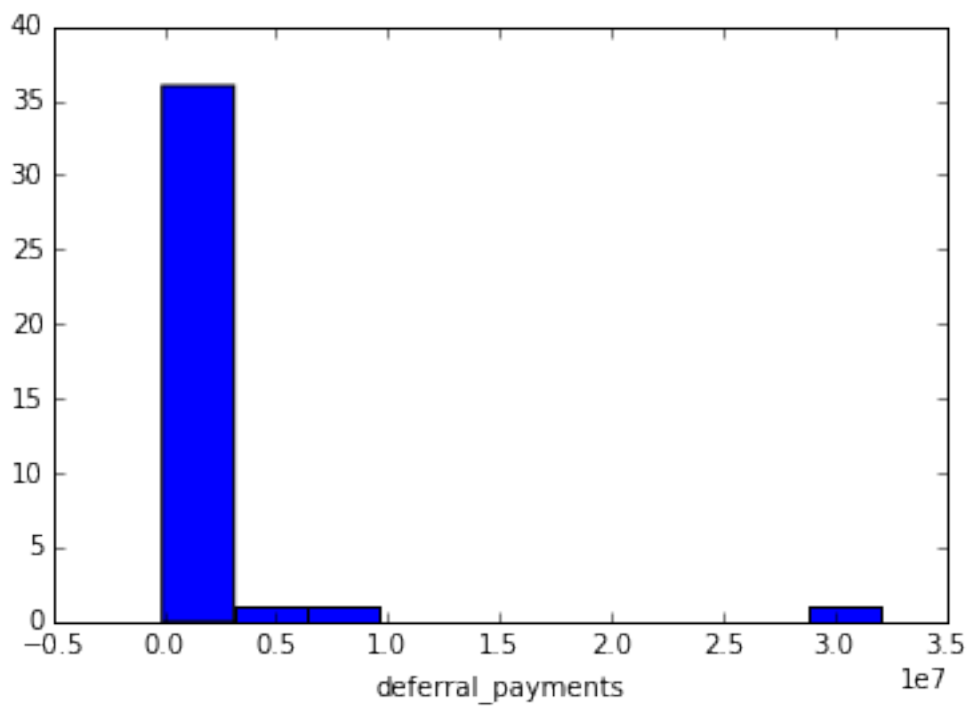
Possible Outlier for salary : TOTAL




Possible Outlier for to\_messages : SHAPIRO RICHARD S




Possible Outlier for deferral\_payments : TOTAL




Possible Outlier for total.payments : TOTAL




Possible Outlier for exercised.stock.options : TOTAL




Possible Outlier for bonus : TOTAL




Possible Outlier for restricted\_stock : TOTAL




Possible Outlier for shared\_receipt\_with\_poi : BELDEN TIMOTHY N




Possible Outlier for restricted\_stock\_deferred : BHATNAGAR SANJAY



Possible Outlier for total\_stock\_value : TOTAL




Possible Outlier for expenses : TOTAL







Possible Outlier for loan\_advances : TOTAL




Possible Outlier for from\_messages : KAMINSKI WINCENTY J




Possible Outlier for other : TOTAL



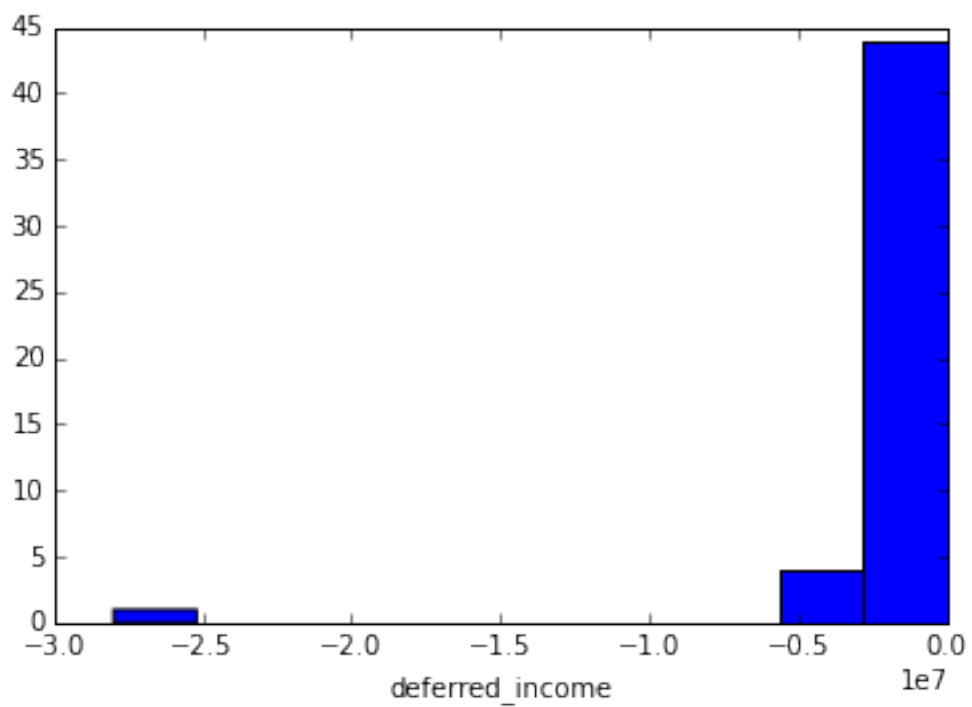
Possible Outlier for from.this.person.to.poi : DELAINEY DAVID W



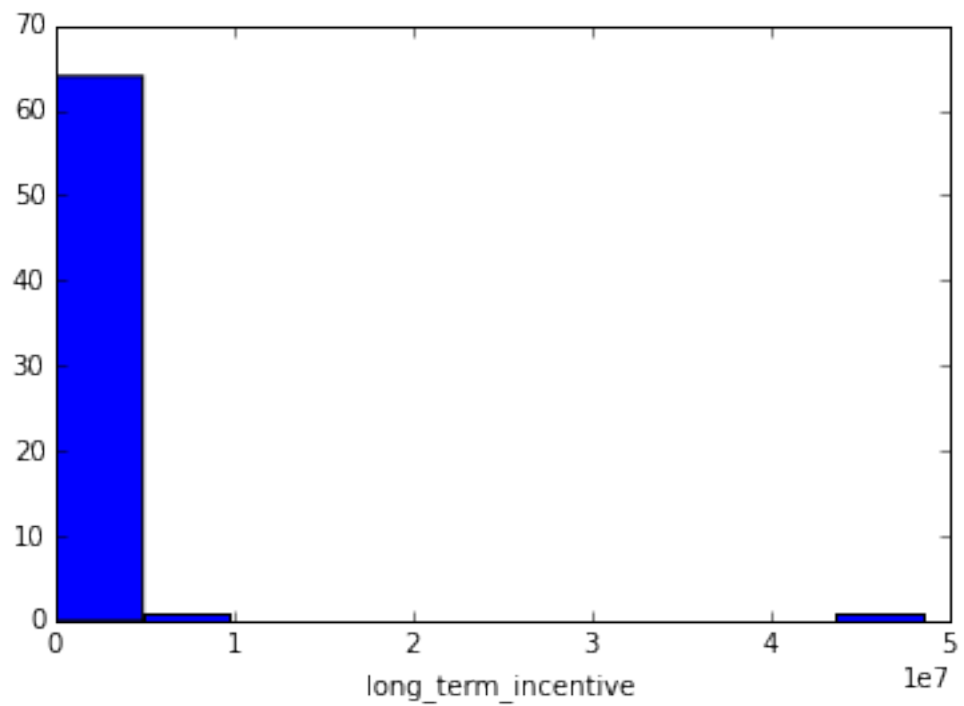
Possible Outlier for director\_fees : TOTAL




Possible Outlier for deferred\_income : BOWEN JR RAYMOND M



Possible Outlier for long\_term\_incentive : TOTAL



Possible Outlier for from\_poi\_to\_this\_person : LAVORATO JOHN J



```

In [15]: # Locate bad data (if any)
         for person in data_dict:
             delete = True
             for feature in data_dict[person]:
                 if feature == "email_address" or feature == 'poi':
                     continue
                 if data_dict[person][feature] == 'NaN':
                     continue
                 if data_dict[person][feature] > 0:
                     delete = False #If any field is positive, then keep the record...
             if delete:
                 print person
                 for feature in data_dict[person]:
                     print feature, data_dict[person][feature] # something is wrong with the data.

```

```

LOCKHART EUGENE E
salary NaN
to_messages NaN
deferral_payments NaN
total_payments NaN
exercised_stock_options NaN
bonus NaN
restricted_stock NaN
shared_receipt_with_poi NaN
restricted_stock_deferred NaN
total_stock_value NaN
expenses NaN
loan_advances NaN
from_messages NaN
other NaN
from_this_person_to_poi NaN
poi False
director_fees NaN
deferred_income NaN
long_term_incentive NaN
email_address NaN
from_poi_to_this_person NaN

```

```

In [16]: # By looking at the histograms and the "person" with the max value,
         # it becomes obvious that "TOTAL" is a sum of the data and not a valid person.
         # Remove TOTAL from the data. Additionally, by since there are only 146 total
         # records, a visual inspection of the names reveals not only the "TOTAL" person,
         # but also a person named "THE TRAVEL AGENCY IN THE PARK". Upon reviewing the
         # accompanying pdf from FindLaw, we learn that "Payments were made by Enron
         # employees on account of business-related travel" to this "person". Since this
         # is not a POI and not a real person, I decided it should also be excluded.
         # Additionally, all the values for "LOCKHART EUGENE E" are NaN, so I excluded this
         # record as well.
         data_dict.pop('TOTAL', 0 )
         data_dict.pop('THE TRAVEL AGENCY IN THE PARK', 0 )
         data_dict.pop('LOCKHART EUGENE E', 0 )

```


```

# Reproduce the histograms and look for possibly other outliers...
for feature in data_dict[data_dict.keys()[0]].keys():
    if feature == "email_address" or feature == 'poi':
        continue
    data = featureFormat(data_dict, [feature])
    for person in data_dict:
        if data_dict[person][feature] == max(data):
            print "Possible Outlier for", feature, ":", person
matplotlib.pyplot.hist(data, bins=10)
matplotlib.pyplot.xlabel(feature)
matplotlib.pyplot.show()


###
# After reviewing the new histograms, while some people appear more than others
# with large values, these are persons of interest and should remain in the dataset.
# No further cleaning is required.

```


Possible Outlier for salary : SKILLING JEFFREY K




Possible Outlier for to\_messages : SHAPIRO RICHARD S




Possible Outlier for deferral\_payments : FREVERT MARK A



Possible Outlier for total\_payments : LAY KENNETH L




Possible Outlier for exercised\_stock\_options : LAY KENNETH L







Possible Outlier for bonus : LAVORATO JOHN J



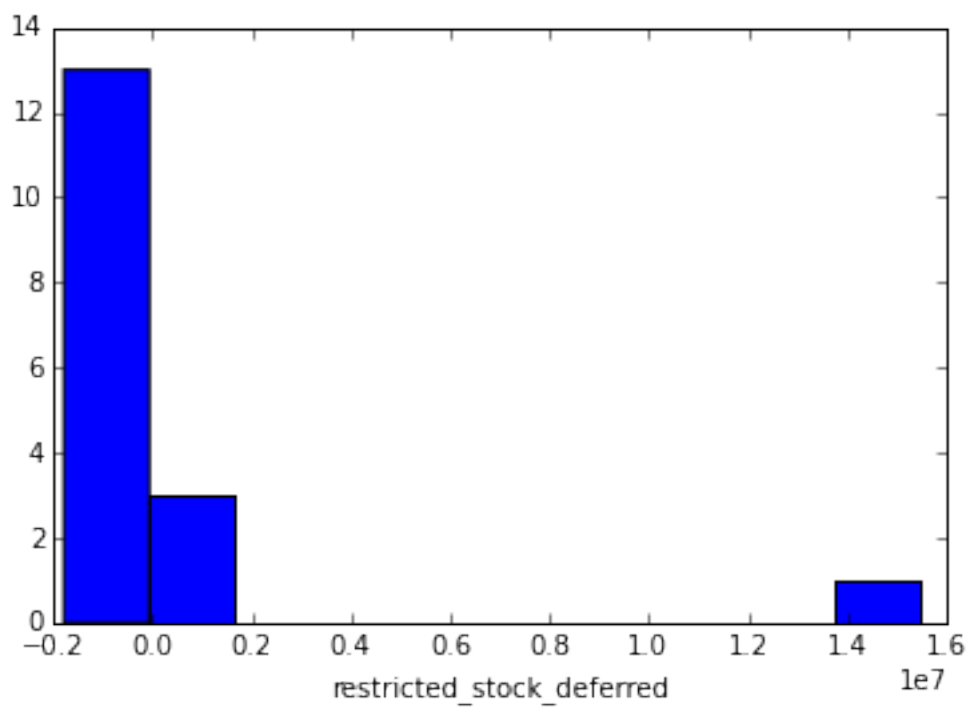
Possible Outlier for restricted\_stock : LAY KENNETH L




Possible Outlier for shared\_receipt\_with\_poi : BELDEN TIMOTHY N




Possible Outlier for restricted\_stock\_deferred : BHATNAGAR SANJAY




Possible Outlier for total\_stock\_value : LAY KENNETH L




Possible Outlier for expenses : MCCLELLAN GEORGE




Possible Outlier for loan\_advances : LAY KENNETH L




Possible Outlier for from\_messages : KAMINSKI WINCENTY J




Possible Outlier for other : LAY KENNETH L




Possible Outlier for from.this.person.to.poi : DELAINEY DAVID W




Possible Outlier for director\_fees : BHATNAGAR SANJAY




Possible Outlier for deferred\_income : BOWEN JR RAYMOND M



Possible Outlier for long\_term\_incentive : MARTIN AMANDA K



Possible Outlier for from\_poi\_to\_this\_person : LAVORATO JOHN J



```

In [17]: ### Task 3: Create new feature(s)
         ### Store to my_dataset for easy export below.
         my_dataset = data_dict

         # Add ratio of bonus to pay
         for person in my_dataset:
             if my_dataset[person]['salary'] > 0 and \
                my_dataset[person]['salary'] != 'NaN' and \
                my_dataset[person]['bonus'] != 'NaN':
                 my_dataset[person]['bonusratio'] = 1.0 * \
                    my_dataset[person]['bonus'] / my_dataset[person]['salary']
             else:
                 my_dataset[person]['bonusratio'] = 'NaN'

         # Add ratio of poi email messages (higher percentage of email
         # to/from a poi could indicate also a poi)
         for person in my_dataset:
             if my_dataset[person]['to_messages'] != 'NaN' and \
                my_dataset[person]['from_messages'] != 'NaN' and \
                my_dataset[person]['from_this_person_to_poi'] != 'NaN' and \
                my_dataset[person]['from_poi_to_this_person'] != 'NaN':
                 poi_email = my_dataset[person]['from_this_person_to_poi'] + \
                    my_dataset[person]['from_poi_to_this_person']
                 all_email = my_dataset[person]['to_messages'] + my_dataset[person]['from_messages']
                 my_dataset[person]['poi_email_ratio'] = 1.0 * poi_email / all_email
             else:
                 my_dataset[person]['poi_email_ratio'] = 'NaN'

In [18]: ### Task 4: Try a variety of classifiers
         ### Please name your classifier clf for easy export below.
         ### Note that if you want to do PCA or other multi-stage operations,
         ### you'll need to use Pipelines. For more info:
         ### http://scikit-learn.org/stable/modules/pipeline.html

         # Below are the classifiers that were tried. Pipelines were used in order to easily setup
         # the flow. A MinMaxScaler was used in the pipeline because some pipelines included PCA,
         # which has been results when used on scaled values.
         # 1: MinMaxScaler, SelectKBest(score_func=f_classif), PCA, GaussianNB
         # {'kbest__k': [5, 10, 15, 20], 'pca__n_components': [2, 3, 4],
         #      'pca__whiten': [True, False]}
         # fit time: 83.653 s
         # Best Params: {'kbest__k': 10, 'pca__n_components': 4, 'pca__whiten': True}
         # Accuracy: 0.83227, Precision: 0.34877, Recall: 0.29750, F1: 0.32110, F2: 0.30651
         #
         # 2: SelectKBest(score_func=f_classif), GaussianNB
         # {'kbest__k': [5, 10, 15, 20]}
         # pipeline 1 fit time: 10.282 s
         # Best Params: {'kbest__k': 5}
         # Accuracy: 0.84653, Precision: 0.41322, Recall: 0.35950, F1: 0.38449, F2: 0.36910
         #
         # 3: MinMaxScaler, SelectKBest(score_func=f_classif), PCA, LogisticRegression

```



```

# {'kbest__k': [5, 10, 15, 20], 'pca__n_components': [2, 3, 4],
#   'pca__whiten': [True, False], 'clf__C': [1, 10, 100, 1000],
#   'clf__solver': ['liblinear', 'newton-cg']}
# fit time: 1083.793 s
# Best Params: {'pca__n_components': 4, 'kbest__k': 10, 'clf__C': 100,
#   'clf__solver': 'liblinear', 'pca__whiten': True}
# Accuracy: 0.85933, Precision: 0.42188, Recall: 0.14850, F1: 0.21967, F2: 0.17061
#
# 4: MinMaxScaler, SelectKBest(score_func=f_classif), PCA, SVC(random_state =42)
# {'kbest__k': [5, 10, 15, 20], 'pca__n_components': [2, 3, 4],
#   'pca__whiten': [True, False], 'clf__C': [1, 10, 100, 1000],
#   'clf__kernel': ['rbf', 'linear']}
# fit time: 5934.505 s
# Best Params: {'kbest__k': 15, 'clf__C': 1000, 'pca__n_components': 4,
#   'pca__whiten': True, 'clf__kernel': 'rbf'}
# Accuracy: 0.81840, Precision: 0.31146, Recall: 0.29900, F1: 0.30510, F2: 0.30141
#
# 5: MinMaxScaler, SelectKBest(score_func=f_classif), PCA,
#   DecisionTreeClassifier(random_state =42)
# {'kbest__k': [5, 10, 15, 20], 'pca__n_components': [2, 3, 4],
#   'pca__whiten': [True, False], 'clf__criterion': ['gini', 'entropy'],
#   'clf__max_depth': [10, 100, 1000, 10000]}
# fit time: 711.177 s
# Best Params: {'pca__n_components': 3, 'clf__criterion': 'entropy',
#   'clf__max_depth': 100, 'kbest__k': 10, 'pca__whiten': True}
# Accuracy: 0.79760, Precision: 0.23161, Recall: 0.22350, F1: 0.22748, F2: 0.22508
#
# 6: MinMaxScaler, SelectKBest(score_func=f_classif), GaussianNB
# {'kbest__k': [5, 10, 15, 20]}
# fit time: 11.212 s
# Best Params: {'kbest__k': 5}
# Accuracy: 0.84653, Precision: 0.41322, Recall: 0.35950, F1: 0.38449, F2: 0.36910
#
# 7: MinMaxScaler, SelectKBest(score_func=f_classif), LogisticRegression
# {'kbest__k': [5, 10, 15, 20], 'clf__C': [1, 10, 100, 1000],
#   'clf__solver': ['liblinear', 'newton-cg']}
# fit time: 213.16 s
# Best Params: {'kbest__k': 20, 'clf__C': 1000, 'clf__solver': 'liblinear'}
# Accuracy: 0.81933, Precision: 0.21371, Recall: 0.13250, F1: 0.16358, F2: 0.14340
#
# 8: MinMaxScaler, SelectKBest(score_func=f_classif),
#   DecisionTreeClassifier(random_state =42)
# {'kbest__k': [5, 10, 15, 20], 'clf__criterion': ['gini', 'entropy'],
#   'clf__max_depth': [10, 100, 1000, 10000]}
# fit time: 104.541 s
# Best Params: {'clf__criterion': 'entropy', 'clf__max_depth': 100, 'kbest__k': 15}
# Accuracy: 0.80153, Precision: 0.23494, Recall: 0.21650, F1: 0.22534, F2: 0.21995
#
# 9: MinMaxScaler, SelectKBest(score_func=f_classif), LogisticRegression
# {'kbest__k': [5, 10, 15, 20], 'clf__C': [1000, 10000, 100000, 1000000],
#   'clf__solver': ['liblinear', 'newton-cg']}
# fit time: 403.069 s
# Best Params: {'kbest__k': 20, 'clf__C': 1000000, 'clf__solver': 'newton-cg'}
# Accuracy: 0.81647, Precision: 0.26336, Recall: 0.20950, F1: 0.23336, F2: 0.21843

```

```

#
# 10: MinMaxScaler, SelectKBest(score_func=f_classif), PCA, SVC(random_state =42)
# {'kbest__k': [5, 10, 15, 20], 'clf__C': [1000, 10000, 100000, 1000000],
#   'clf__kernel': ['rbf', 'linear']}
# fit time: 42313.693 s
# Best Params: {'kbest__k': 20, 'clf__C': 100000, 'clf__kernel': 'rbf'}
# Accuracy: 0.81113, Precision: 0.29018, Recall: 0.28800, F1: 0.28908, F2: 0.28843
#
# 11: MinMaxScaler, SelectKBest(score_func=f_classif), PCA,
#     RandomForestClassifier(random_state =42)
# {'kbest__k': [5, 10, 15, 20], 'clf__criterion': ['gini', 'entropy'],
#   'clf__n_estimators': [10, 100, 1000, 10000]}
# fit time: 66037.228 s
# Best Params: {'clf__criterion': 'gini', 'kbest__k': 5, 'clf__n_estimators': 100}
# Accuracy: 0.85973, Precision: 0.44503, Recall: 0.21050, F1: 0.28581, F2: 0.23530

# The best classifier turned out to be the GaussianNB when run with
# SelectKBest(k=5, score_func=f_classif). The pipeline fit was about 10 seconds. Adding in
# scaling made no change to the accuracy or precision, which was expected.

# Setup the pipeline
# GaussianNB gave the best results with no scaler and no PCA. This is the classifier that
# will be used to generate the project pickle files. All the above tests were performed in
# the same manner by modifying the pipeline.
pipeline = Pipeline(steps=[('kbest', SelectKBest(score_func=f_classif)),
                           ('clf', GaussianNB())
                           ])
p_params = {'kbest__k': [5, 10, 15, 20]}

In [19]: ### Task 5: Tune your classifier to achieve better than .3 precision and recall
### using our testing script. Check the tester.py script in the final project
### folder for details on the evaluation method, especially the test_classifier
### function. Because of the small size of the dataset, the script uses
### stratified shuffle split cross validation. For more info:
### http://scikit-learn.org/stable/modules/generated/sklearn.cross\_validation.StratifiedShuffleSplit

# Example starting point. Try investigating other evaluation techniques!
# Create a feature_list that has all features on it. We will use selectKBest to determine the
all_features = data_dict[my_dataset.keys()[0]].keys()
all_features.remove('poi')
all_features.remove('email_address')
features_list = ['poi']
features_list.extend(all_features)

### Extract features and labels from dataset for local testing
data = featureFormat(my_dataset, features_list, sort_keys = True)
labels, features = targetFeatureSplit(data)

# Create 1000 random test sets to go over...
shuffle = StratifiedShuffleSplit(labels, n_iter=1000, test_size=0.3, random_state=42)

#Select a scoring function. Turns out recall is not so good...
#scorer = 'recall'
scorer = 'f1'

```

```

In [20]: # Find the best classifier
         t0 = time()
         p_grid = GridSearchCV(pipeline, param_grid=p_params, cv=shuffle,
                                scoring=scorer, verbose=0)
         p_grid.fit(features, labels)
         print "fit time:", round(time()-t0, 3), "s"

fit time: 9.98 s

In [21]: # Results....
         features_selected = p_grid.best_estimator_.named_steps['kbest'].get_support()
         x=0
         print "Feature(Score) Selected:"
         for feat in features_list:
             if feat == "poi":
                 continue
             if features_selected[x] == True:
                 print '    %s(%f)' % (feat, p_grid.best_estimator_.named_steps['kbest'].scores_[x])
             x += 1
         print "Score:", p_grid.best_score_
         print "Params:", p_grid.best_params_

         clf = p_grid.best_estimator_
         test_classifier(clf, my_dataset, features_list)

Feature(Score) Selected:
    deferred_income(11.458477)
    bonus(20.792252)
    total_stock_value(24.182899)
    salary(18.289684)
    exercised_stock_options(24.815080)
Score: 0.309143105476
Params: {'kbest_k': 5}
Pipeline(steps=[('kbest', SelectKBest(k=5, score_func=<function f_classif at 0x000000001AD2F668>)), ('clf',
    Accuracy: 0.84653      Precision: 0.41322      Recall: 0.35950      F1: 0.38449      F2: 0.35950
    Total predictions: 15000      True positives: 719      False positives: 1021      False negatives: 7890)]),
In [22]: ### Task 6: Dump your classifier, dataset, and features_list so anyone can
         ### check your results. You do not need to change anything below, but make sure
         ### that the version of poi_id.py that you submit can be run on its own and
         ### generates the necessary .pkl files for validating your results.

         dump_classifier_and_data(clf, my_dataset, features_list)

```